input_ids

[B, S]

model · Qwen3Model

≈11.4B

input_ids

[B, S]

Embedding

embed_tokens · 151936 → 4096 · 622M

[B, S, 4096]

layers · Qwen3DecoderLayer

× 32

10.8B

hidden_states

Qwen3RMSNorm

input_layernorm · 4096 · 4.10K

Qwen3Attention

self_attn · 67.1M

+

Qwen3RMSNorm

post_attention_layernorm · 4096 · 4.10K

Qwen3MLP

mlp · 273M

+

hidden_states

Qwen3RMSNorm

norm · 4096 · 4.10K

last_hidden_state

[B, S, 4096]

Linear

lm_head · 4096 → 151936 · 622M

[B, S, 151936]

logits